

Sangdon Park

Assistant Professor, POSTECH GSAI/CSE

 sangdon.github.io
 sangdon@postech.ac.kr
 github.com/sangdon
 Google Scholar

RESEARCH INTERESTS

My research focuses on the foundation of trustworthy machine learning for safety-critical real-world AI, including agentic and physical AI by developing principled learning algorithms for trustworthy adaptation in the real world.

EDUCATION

University of Pennsylvania

Ph.D. in Computer and Information Science

Philadelphia, USA

2021

- Advisors: [Insup Lee](#) and [Osbert Bastani](#)
- Thesis: *Uncertainty Estimation Toward Safe AI*
- Committee: [Kostas Daniilidis](#), [Nikolai Matni](#), [Edgar Dobriban](#), and [Kilian Q. Weinberger](#)

Seoul National University

M.S. in Electrical and Computer Engineering

Seoul, Korea

2012

- Advisor: [Kyoung Mu Lee](#)
- Thesis: *Abnormal Object Detection by Transformed-Canonical Scene Generation*

Seoul National University

B.S. in Computer Science and Engineering

Seoul, Korea

2010

- Thesis Advisor: [Byoung-Tak Zhang](#)
- Thesis: *Behavioral Intelligence for Crowd Avatar in 3D Virtual Worlds*

EMPLOYMENT

POSTECH

Assistant Professor

Pohang, Korea

Aug. 2023–Now

Georgia Institute of Technology

Postdoctoral Researcher (Mentor: [Taesoo Kim](#))

Atlanta, USA

Sep. 2021–Jul. 2023

Google Cloud AI

Research Intern (Host: [Kihyuk Sohn](#))

Sunnyvale, USA

Summer 2020

Biointelligence Laboratory, Seoul National University

Undergraduate Researcher

Seoul, Korea

2008–2010

Republic of Korea Army

Military Service

Korea

2006–2008

PUBLICATIONS

- [1] M. Lee*, Y. Jung*, and **S. Park**, “Online Conformal Abstention Under Adversarial Bandit Feedback”, in The Fortieth Annual Conference on Neural Information Processing Systems (**NeurIPS**), 2026.
- [2] J. Park, N. Park, S. Lee, Y.-C. Jhi, J. Cho, and **S. Park**, “MJ: Multi-turn LLM Jailbreaking via Decomposed Credit Assignment”, in The Fortieth Annual Conference on Neural Information Processing Systems (**NeurIPS**), 2026.
- [3] M. Park*, S. Lee*, K. Kim, J. Heo, S. Cho, S. Li, **S. Park**, and D. Kim, “Transductive Generalization for GNNs via Optimal Transport”, in The Fortieth Annual Conference on Neural Information Processing Systems (**NeurIPS**), 2026.
- [4] J. Choi*, J. Yang*, and **S. Park**, “SAUL: Sharpness-Aware Augmented-Lagrangian Unlearning”, in The 2026 Conference on Empirical Methods in Natural Language Processing (**EMNLP Findings**), 2026.
- [5] K. Kim*, Y. Choi*, S. Lee, S. Jun, D. Kim†, and S. Park†, “The Interplay of Harness Design and Post-Training in LLM Agents”, in The 2026 Conference on Empirical Methods in Natural Language Processing (**EMNLP Findings**), 2026.
- [6] S. Lee, H. Kim, and **S. Park**, “Can We Stop Malicious AI? KILLBENCH: A Benchmark for External AI Kill Switch Feasibility”, in The 2026 Conference on Empirical Methods in Natural Language Processing (**EMNLP Findings**), 2026.
- [7] X. Cheng, F. Sang, Y. Yang, H. Zhang, **S. Park**, X. Zhang, and T. Kim, “The Illusion of Rust Safety: Detecting Modular Unsafe Functions with LLMs”, in The ACM Conference on Computer and Communications Security (**CCS**), 2026.
- [8] J. Jeong, T. Kim, and **S. Park**, “Towards Functional Correctness of Large Code Models with Selective Generation”, in International Conference on Machine Learning (**ICML**), 2026.
- [9] J. Hwang, **S. Park**, and J. Ok, “LLM Watermark Evasion via Bias Inversion”, in International Conference on Machine Learning (**ICML**), 2026.
- [10] X. Cheng*, S. Park*, H. Han, X. Zhang, and T. Kim, “Ruby: Unmasking Unsafe Rust in Stripped Binaries via Machine Learning”, in International Conference on Dependable Systems and Networks (**DSN**), 2026.
- [11] J. Yang*, K. Kim*, and **S. Park**, “Online Conformal Prediction with Adversarial Semi-bandit Feedback via Regret Minimization”, in International Conference on Learning Representations (**ICLR**), 2026.
- [12] T. Kim, H. Han, S. Park, D. R. Jeong, D. Kim, D. Kim, E. Kim, J. Kim, J. Wang, K. Kim, S. Ji, W. Song, H. Zhao, A. Chin, G. Lee, K. Stevens, M. Alharthi, Y. Zhai, C. Zhang, J. Jang, Y. Jang, A. Askar, D. Kim, F. Fleischer, J. Cho, J. Kim, K. Ko, I. Yun, **S. Park**, D. Baik, H. Lee, H. Heo, M. Gwon, M. Lee, M. Baek, S. Min, W. Kim, Y. Jin, Y. Park, Y. Choi, J. Jung, G. Lee, J. Jang, K. Kim, Y. Cha, and Y. Kim, *ATLANTIS: AI-driven Threat Localization, Analysis, and Triage Intelligence System*, 2025. arXiv: [2509.14589](https://arxiv.org/abs/2509.14589) [[cs](#).[CR](#)].
- [13] K. Kim, Y. Choi, H. Kim, D. Kim, and **S. Park**, “ChronoBias: A Benchmark for Evaluating Time-conditional Group Bias in the Time-sensitive Knowledge of Large Language Models”, in Findings of the Association for Computational Linguistics: EMNLP (**EMNLP Findings**), 2025.
- [14] J. Hwang, J. Park, H. Park, D. Kim, **S. Park**, and J. Ok, “Retrieval-Augmented Generation with Estimation of Source Reliability”, in Empirical Methods in Natural Language Processing (**EMNLP**), 2025.
- [15] S. Moon*, M. Lee*, **S. Park**, and D. Kim, “Holistic Unlearning Benchmark: A Multi-Faceted Evaluation for Text-to-Image Diffusion Model Unlearning”, in International Conference on Computer Vision (**ICCV**), 2025.

- [16] H. Wang, Z. Yang, **S. Park**, Y. Yang, S. Kim, W. Lunardi, M. Andreoni, T. Kim, and W. Lee, “SOUNDBOOST: Effective RCA and Attack Detection for UAV via Acoustic Side-Channel”, in International Conference on Dependable Systems and Networks (**DSN**), 2025.
- [17] M. Lee*, K. Kim*, T. Kim, and **S. Park**, “Selective Generation for Controllable Language Models”, in Neural Information Processing Systems (**NeurIPS**), 2024.
- [18] S. Li, **S. Park**, I. Lee, and O. Bastani, “TRAQ: Trustworthy Retrieval Augmented Question Answering via Conformal Prediction”, in North American Chapter of the Association for Computational Linguistics (**NAACL**), 2024.
- [19] H. Park*, J. Hwang*, S. Mun, **S. Park**, and J. Ok, “MedBN: Robust Test Time Adaptation against Malicious Test Samples”, in Computer Vision and Pattern Recognition (**CVPR**), 2024.
- [20] W. Si, **S. Park**, I. Lee, E. Dobriban, and O. Bastani, “PAC Prediction Sets Under Label Shift”, in International Conference on Learning Representations (**ICLR**), 2024.
- [21] W. Si, S. Li, **S. Park**, I. Lee, and O. Bastani, “Angelic Patches for Improving Third-Party Object Detector Performance”, in Computer Vision and Pattern Recognition (**CVPR**), 2023.
- [22] **S. Park**, O. Bastani, and T. Kim, “ACon²: Adaptive Conformal Consensus for Provable Blockchain Oracles”, in USENIX Security Symposium (**Security**), 2023.
- [23] R. Kaur, K. Sridhar, **S. Park**, Y. Yang, S. Jha, A. Roy, O. Sokolsky, and I. Lee, “CODiT: Conformal Out-of-Distribution Detection in Time-Series Data for Cyber-Physical Systems”, in International Conference on Cyber-Physical Systems (**ICCPs**), 2023.
- [24] **S. Park**, E. Dobriban, I. Lee, and O. Bastani, “PAC Prediction Sets for Meta-Learning”, in Neural Information Processing Systems (**NeurIPS**), 2022.
- [25] S. Li, **S. Park**, X. Ji, I. Lee, and O. Bastani, *Towards PAC Multi-Object Detection and Tracking*, 2022. arXiv: 2204.07482 [cs.CV].
- [26] S. Jang, **S. Park**, I. Lee, and O. Bastani, “Sequential Covariate Shift Detection Using Classifier Two-Sample Tests”, in International Conference on Machine Learning (**ICML**), 2022.
- [27] **S. Park**, E. Dobriban, I. Lee, and O. Bastani, “PAC Prediction Sets Under Covariate Shift”, in International Conference on Learning Representations (**ICLR**), 2022.
- [28] R. Kaur, S. Jha, A. Roy, **S. Park**, E. Dobriban, O. Sokolsky, and I. Lee, “iDECODE: In-distribution Equivariance for Conformal Out-of-distribution Detection”, in Association for the Advancement of Artificial Intelligence (**AAAI**), 2021.
- [29] **S. Park**, S. Li, I. Lee, and O. Bastani, “PAC Confidence Predictions for Deep Neural Network Classifiers”, in International Conference on Learning Representations (**ICLR**), 2021.
- [30] **S. Park**, O. Bastani, J. Weimer, and I. Lee, “Calibrated prediction with covariate shift via unsupervised domain adaptation”, in International Conference on Artificial Intelligence and Statistics (**AISTATS**), 2020.
- [31] **S. Park**, O. Bastani, N. Matni, and I. Lee, “PAC Confidence Sets for Deep Neural Networks via Calibrated Prediction”, in International Conference on Learning Representations (**ICLR**), 2020.
- [32] **S. Park**, J. Weimer, and I. Lee, “Resilient linear classification: an approach to deal with attacks on training data”, in International Conference on Cyber-Physical Systems (**ICCPs**), 2017.
- [33] J. Oh, T. M. Howard, M. R. Walter, D. Barber, M. Zhu, **S. Park**, A. Suppe, L. Navarro-Serment, F. Duvallet, A. Boularias, et al., “Integrated intelligence for human-robot teams”, in International Symposium on Experimental Robotics (**ISER**), 2016.
- [34] **S. Park**, W. Kim, and K. M. Lee, “Abnormal object detection by canonical scene-based contextual model”, in European Conference on Computer Vision (**ECCV**), 2012.

SCHOLARSHIPS AND AWARDS

- Best Paper Award from CKAIA (한국인공지능학회) 2026
- DARPA AIxCC Winner as [Team Atlanta](#) (\$4M Team Prize) 2025
- Best Paper Finalist from CKAIA (한국인공지능학회) 2025
- Bang Seung-Yang Graduate Fellowship from POSTECH CSE (awardee: Jaewoo Jeong) 2025
- POSTECH GSAI BK21 Best Paper Award 2025
- NeurIPS'24 Spotlight Paper (top 2.08%) 2024
- DARPA AIxCC Finalist as [Team Atlanta](#) (\$2M Team Prize) 2024
- NeurIPS'23 Outstanding Reviewer Award 2023
- ICML'23 TEACH Workshop Best Paper Award 2023
- ICCPS'23 Best Paper Award Finalist 2023
- NeurIPS'21 Outstanding Reviewer Award (top 8% of reviewers) 2021
- Korea Cyber-Security Paper Competition Best Paper Award (\$4,500) 2018
- PhD fellowship at University of Pennsylvania 2013–2021
- Distinguished MS Dissertation Award at Seoul National University 2012
- Academic Performance Scholarship 2009
- National Science and Engineering Undergraduate Scholarship 2003– 2008

SERVICE

- **Area Chair**
NeurIPS'24–26, ICML'25–26, ICLR'26–27
- **Reviewer**
NeurIPS'21–23, ICML'21–23, ICLR'22–24, AAAI'25, Journal of the Royal Statistical Society: Series B
- **External Reviewer**
S&P'21, S&P'22, Security'22, Security'23, Security'24, NDSS'24

TEACHING

- **Instructor** at POSTECH
ML (CSED515) Fall 2026
- **Instructor** at POSTECH
AI Security (CSED490H) Spring 2026
- **Instructor** at POSTECH
ML (CSED515) Fall 2025
- **Instructor** at POSTECH
Trustworthy ML (AIGS703L / CSED703L) Spring 2025
- **Instructor** at POSTECH
Trustworthy ML (AIGS703L / CSED703L) Fall 2024
- **Instructor** at POSTECH
Discrete Mathematics (CSED261) Spring 2024
- **Instructor** at POSTECH
Trustworthy ML (AIGS703L / CSED703L) Fall 2023

- **Teaching Assistant** at University of Pennsylvania
Machine Perception (CIS580) Spring 2015
- **Teaching Assistant** at University of Pennsylvania
Computer Vision and Computational Photography (CIS581) Fall 2014
- **Teaching Assistant** at Seoul National University
Linear Algebra for Electrical Systems Fall 2010
- **Instructor** at Seoul National University
1st Free Computer Education for Gwanak-gu Community Youth Feb. 2008

TALKS

- **From LLMs to Agentic and Physical AI**
POSTECH x 한국무역협회 최고경영자과정 Sep. 2026
- **Towards Agentic AI Alignment**
데이터보안프라이버시연구회 Jun. 2026
- **AI Red Teaming: Towards AI Alignment**
KCSE (한국소프트웨어공학학술대회) Feb. 2026
- **AI Alignment**
Nxt AI Nov. 2025
KAIST GSIS Nov. 2025
- **AI Red Teaming and Agentic AI Security**
OSSTEM Nov. 2025
Samsung SDS Nov. 2025
WITHCON Nov. 2025
- **AI System for Discovering and Patching Vulnerable Code by Team Atlanta**
POSCO DX Oct. 2025
- **Introduction to Trustworthy LLMs**
KISC (한국통신학회 단기강좌) Jul. 2025
- **Trustworthy LLMs**
Best of the Best (BoB) Jul. 2025
- **Selective Generation Towards Trustworthy LLMs**
NYU Frontier Lab Jul. 2025
KIISC AI Security Group (정보보호학회 AI보안 연구회) Jul. 2025
- **Toward Trustworthy Large Language Models**
PKNU AI Dec. 2024
POSTECH AI Day Dec. 2024
- **Rethinking and Harnessing Trustworthiness of Generative AI for Security**
Samsung Security Tech Forum (SSTF) – *Invited Talk* Sep. 2024
- **Trustworthy Military-AI: AI Controllability**
REAIM Sep. 2024
- **Trend on Trustworthy Language Models**
National Statistics Development Forum Sep. 2024
- **Trustworthy AI: A Compositional Perspective**
POSTECH AI/CSE Apr. 2024
- **Conformal Prediction for Trustworthy AI**
Korean AI Association Winter School Feb. 2024

- **Uncertainty Learning for Trustworthy and Secure AI**
 - POSTECH AI/CSE Mar. 2023
 - KAIST EE Apr. 2023
 - SNU CSE Apr. 2023
 - Korea University CSE Jul. 2023
 - SNU IPAI Jul. 2023
 - SNU Frontier Summer School Aug. 2023
 - UNIST IB Oct. 2023
 - KAIST EE Jan. 2024
 - CAU AI May 2024
 - Korea GSS Jun. 2024
- **PAC Prediction Sets for AI Safety**
 - ICML Workshop DFUQ 2022 – *Invited Talk* Jul. 2022
- **Uncertainty Quantification via PAC Prediction Sets**
 - DGIST Dec. 2021
- **From Verification to Learning for Defense against Adversarial Examples in Neural Networks**
 - KAIST GSIS Aug. 2018
 - Hanyang University Aug. 2018
 - KIISC Aug. 2018